



Regularization
Feature scaling

Regularization

- Before we dive straight into coding regularization with Scikit-Learn, we need to discuss a few more relevant topics:
 - Feature Scaling
 - Cross Validation

Regularization

- Before we dive straight into coding regularization with Scikit-Learn, we need to discuss a few more relevant topics:
 - Feature Scaling
 - Cross Validation

Regularization – Feature Scaling

- Feature scaling provides many benefits to our machine learning process!
- Some machine learning models that rely on distance metrics (e.g. KNN) **require** scaling to perform well.
- Let's discuss the main ideas behind feature scaling...

Regularization – Feature Scaling

- Feature scaling improves the convergence of steepest descent algorithms, which do not possess the property of scale invariance.
- If features are on different scales, certain weights may update faster than others since the feature values x_j play a role in the weight updates.

Regularization – Feature Scaling

- Critical benefit of feature scaling related to gradient descent.
- There are some ML Algos where scaling won't have an effect (e.g. CART based methods).

Regularization – Feature Scaling

- Scaling the features so that their respective ranges are uniform is important in comparing measurements that have different units.
- Allows us directly compare model coefficients to each other.

Regularization – Feature Scaling

- Feature scaling caveats:
 - Must always scale new unseen data before feeding to model.
 - Effects direct interpretability of feature coefficients
 - Easier to compare coefficients to one another, harder to relate back to original unscaled feature.

Regularization – Feature Scaling

- Feature scaling benefits:
 - Can lead to great increases in performance.
 - Absolutely necessary for some models.
 - Virtually no “real” downside to scaling features.

Regularization – Feature Scaling

- Two main ways to scale features:
 - Standardization
 - Rescales data to have a mean (μ) of 0 and standard deviation (σ) of 1.
 - Normalization
 - Rescales all data values to be between 0-1.

Regularization – Feature Scaling-standardization

- Standardization:
 - Rescales data to have a mean (μ) of 0 and standard deviation (σ) of 1 (unit variance).

$$X_{changed} = \frac{X - \mu}{\sigma}$$

Regularization – Feature Scaling-standardization

- Standardization:
 - Namesake can be confusing since this is also referred to as “Z-score normalization”.

$$X_{changed} = \frac{X - \mu}{\sigma}$$

Regularization – Feature Scaling- Normalization

- Normalization:
 - Scales all data values to be between 0 and 1.

$$X_{changed} = \frac{X - X_{min}}{X_{max} - X_{min}}$$

Regularization – Feature Scaling

- There are many more methods of scaling features and Scikit-Learn provides easy to use classes that “fit” and “transform” feature data for scaling.
- Let’s quickly discuss the fit and transform calls in more detail when it comes to scaling.

Regularization – Feature Scaling

- A `.fit()` method call simply calculates the necessary statistics (`Xmin`, `Xmax`, `mean`, `standard deviation`).
- A `.transform()` call actually scales data and returns the new scaled version of data.
- Previously saw a similar process for polynomial feature conversion.

Regularization – Feature Scaling

- Very important consideration for fit and transform:
 - We only **fit** to training data.
 - Calculating statistical information should only come from training data.
 - Don't want to assume prior knowledge of the test set!

Regularization – Feature Scaling

- Using the full data set would cause **data leakage**:
 - Calculating statistics from full data leads to some information of the test set leaking into the training process upon `transform()` conversion.

Regularization – Feature Scaling

- **Feature scaling process:**
 - **Perform train test split**
 - **Fit to training feature data**
 - **Transform training feature data**
 - **Transform test feature data**

Regularization – Feature Scaling

- Do we need to scale the label?
 - In general it is not necessary nor advised.
 - Normalising the output distribution is altering the definition of the target.
 - Predicting a distribution that doesn't mirror your real-world target.
 - Can negatively impact stochastic gradient descent.